

# 3d-image-recognition-using-cnn-1

March 3, 2024

```
[124]: import pandas as pd
import numpy as np
import tensorflow as tf
from tensorflow.keras import datasets, layers, models
import matplotlib.pyplot as plt
```

```
[125]: (X_train, y_train), (X_test, y_test) = datasets.cifar10.load_data()
X_train.shape
```

```
[125]: (50000, 32, 32, 3)
```

```
[126]: X_test.shape
```

```
[126]: (10000, 32, 32, 3)
```

```
[127]: y_train.shape
```

```
[127]: (50000, 1)
```

```
[128]: y_train[:10]
```

```
[128]: array([[6],
        [9],
        [9],
        [4],
        [1],
        [1],
        [2],
        [7],
        [8],
        [3]], dtype=uint8)
```

```
[129]: y_train = y_train.reshape(-1,)
y_train[:10]
```


```
[129]: array([6, 9, 9, 4, 1, 1, 2, 7, 8, 3], dtype=uint8)
```

```
[130]: y_test = y_test.reshape(-1)
```


```
[131]: classes =  
        ↪ ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]
```

```
[132]: def plot_sample(X, y, index):  
        plt.figure(figsize = (15,2))  
        plt.imshow(X[index])  
        plt.xlabel(classes[y[index]])
```

```
[133]: plot_sample(X_train, y_train, 40)
```



```
[134]: plot_sample(X_train, y_train, 1)
```



```
[123]: X_train = X_train / 255  
        X_test = X_test / 255
```

```
[147]: ann = models.Sequential([  
        layers.Flatten(input_shape=(32,32,3)),
```

```

        layers.Dense(3000, activation='relu'),
        layers.Dense(1000, activation='relu'),
        layers.Dense(10, activation='sigmoid')
    ])

ann.compile(optimizer='SGD',
            loss='sparse_categorical_crossentropy',
            metrics=['accuracy'])

ann.fit(X_train, y_train, epochs=5)

```

Epoch 1/5  
 888/1563 [=====>...] - ETA: 25s - loss: nan - accuracy:  
 0.0991

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
Cell In[147], line 12
      1 ann = models.Sequential([
      2     layers.Flatten(input_shape=(32,32,3)),
      3     layers.Dense(3000, activation='relu'),
      4     layers.Dense(1000, activation='relu'),
      5     layers.Dense(10, activation='sigmoid')
      6 ])
      7
      8 ann.compile(optimizer='SGD',
      9             loss='sparse_categorical_crossentropy',
     10             metrics=['accuracy'])
--> 12 ann.fit(X_train, y_train, epochs=5)

File ~\AppData\Roaming\Python\Python311\site-packages\keras\src\utils\traceback_utils.py:65, in filter_traceback.<locals>.error_handler(*args, **kwargs)
     63 filtered_tb = None
     64 try:
--> 65     return fn(*args, **kwargs)
     66 except Exception as e:
     67     filtered_tb = _process_traceback_frames(e.__traceback__)

File ~\AppData\Roaming\Python\Python311\site-packages\keras\src\engine\training.py:1807, in Model.fit(self, x, y, batch_size, epochs, verbose, callbacks, validation_split, validation_data, shuffle, class_weight, sample_weight, initial_epoch, steps_per_epoch, validation_steps, validation_batch_size, validation_freq, max_queue_size, workers, use_multiprocessing)
    1799 with tf.profiler.experimental.Trace(
    1800     "train",
    1801     epoch_num=epoch,
    (...)
    1804     _r=1,

```

```

1805 ):
1806     callbacks.on_train_batch_begin(step)
-> 1807     tmp_logs = self.train_function(iterator)
1808     if data_handler.should_sync:
1809         context.async_wait()

```

File

```

-> ~\AppData\Roaming\Python\Python311\site-packages\tensorflow\python\util\traceback_utils.
py:150, in filter_traceback.<locals>.error_handler(*args, **kwargs)
148 filtered_tb = None
149 try:
--> 150     return fn(*args, **kwargs)
151 except Exception as e:
152     filtered_tb = _process_traceback_frames(e.__traceback__)

```

File

```

-> ~\AppData\Roaming\Python\Python311\site-packages\tensorflow\python\eager\polymorphic_function.
py:832, in Function.__call__(self, *args, **kwargs)
829 compiler = "xla" if self._jit_compile else "nonXla"
831 with OptionalXlaContext(self._jit_compile):
--> 832     result = self._call(*args, **kwargs)
834 new_tracing_count = self.experimental_get_tracing_count()
835 without_tracing = (tracing_count == new_tracing_count)

```

File

```

-> ~\AppData\Roaming\Python\Python311\site-packages\tensorflow\python\eager\polymorphic_function.
py:868, in Function._call(self, *args, **kwargs)
865     self._lock.release()
866     # In this case we have created variables on the first call, so we run
-> the
867     # defunned version which is guaranteed to never create variables.
--> 868     return tracing_compilation.call_function(
869         args, kwargs, self._no_variable_creation_config
870     )
871 elif self._variable_creation_config is not None:
872     # Release the lock early so that multiple threads can perform the call
873     # in parallel.
874     self._lock.release()

```

File

```

-> ~\AppData\Roaming\Python\Python311\site-packages\tensorflow\python\eager\polymorphic_function.
py:139, in call_function(args, kwargs, tracing_options)
137 bound_args = function.function_type.bind(*args, **kwargs)
138 flat_inputs = function.function_type.unpack_inputs(bound_args)
--> 139 return function._call_flat( # pylint: disable=protected-access
140     flat_inputs, captured_inputs=function.captured_inputs
141 )

```

```

File_
↳ ~\AppData\Roaming\Python\Python311\site-packages\tensorflow\python\eager\polymorphic_function.py:1323, in ConcreteFunction._call_flat(self, tensor_inputs, captured_inputs)
    1319 possible_gradient_type = gradients_util.PossibleTapeGradientTypes(args)
    1320 if (possible_gradient_type == gradients_util.POSSIBLE_GRADIENT_TYPES_NONE and
    1321     not executing_eagerly):
    1322     # No tape is watching; skip to running the function.
-> 1323     return self._inference_function.call_preflattened(args)
    1324 forward_backward = self._select_forward_and_backward_functions(
    1325     args,
    1326     possible_gradient_type,
    1327     executing_eagerly)
    1328 forward_function, args_with_tangents = forward_backward.forward()

```

```

File_
↳ ~\AppData\Roaming\Python\Python311\site-packages\tensorflow\python\eager\polymorphic_function.py:216, in AtomicFunction.call_preflattened(self, args)
    214 def call_preflattened(self, args: Sequence[core.Tensor]) -> Any:
    215     """Calls with flattened tensor inputs and returns the structured
    ↳ output."""
-> 216     flat_outputs = self.call_flat(*args)
    217     return self.function_type.pack_output(flat_outputs)

```

```

File_
↳ ~\AppData\Roaming\Python\Python311\site-packages\tensorflow\python\eager\polymorphic_function.py:251, in AtomicFunction.call_flat(self, *args)
    249 with record.stop_recording():
    250     if self._bound_context.executing_eagerly():
-> 251         outputs = self._bound_context.call_function(
    252             self.name,
    253             list(args),
    254             len(self.function_type.flat_outputs),
    255         )
    256     else:
    257         outputs = make_call_op_in_graph(
    258             self,
    259             list(args),
    260             self._bound_context.function_call_options.as_attrs(),
    261         )

```

```

File_
↳ ~\AppData\Roaming\Python\Python311\site-packages\tensorflow\python\eager\context.py:1486, in Context.call_function(self, name, tensor_inputs, num_outputs)
    1484 cancellation_context = cancellation.context()
    1485 if cancellation_context is None:
-> 1486     outputs = execute.execute(
    1487         name.decode("utf-8"),
    1488         num_outputs=num_outputs,
    1489         inputs=tensor_inputs,

```

```

1490     attrs=attrs,
1491     ctx=self,
1492 )
1493 else:
1494     outputs = execute.execute_with_cancellation(
1495         name.decode("utf-8"),
1496         num_outputs=num_outputs,
1497         (...)
1500         cancellation_manager=cancellation_context,
1501     )

```

File

```

→ ~\AppData\Roaming\Python\Python311\site-packages\tensorflow\python\eager\execute.
py:53, in quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
    51 try:
    52     ctx.ensure_initialized()
--> 53     tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_name
    54                                     inputs, attrs, num_outputs)
    55 except core._NotOkStatusException as e:
    56     if name is not None:

```

KeyboardInterrupt:

```

[148]: from sklearn.metrics import confusion_matrix , classification_report
import numpy as np
y_pred = ann.predict(X_test)
y_pred_classes = [np.argmax(element) for element in y_pred]

print("Classification Report: \n", classification_report(y_test,
→ y_pred_classes))

```

313/313 [=====] - 2s 7ms/step

Classification Report:

	precision	recall	f1-score	support
0	0.10	1.00	0.18	1000
1	0.00	0.00	0.00	1000
2	0.00	0.00	0.00	1000
3	0.00	0.00	0.00	1000
4	0.00	0.00	0.00	1000
5	0.00	0.00	0.00	1000
6	0.00	0.00	0.00	1000
7	0.00	0.00	0.00	1000
8	0.00	0.00	0.00	1000
9	0.00	0.00	0.00	1000
accuracy			0.10	10000

macro avg	0.01	0.10	0.02	10000
weighted avg	0.01	0.10	0.02	10000

D:\anaconda\Lib\site-packages\sklearn\metrics\\_classification.py:1469:  
 UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

D:\anaconda\Lib\site-packages\sklearn\metrics\\_classification.py:1469:  
 UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

D:\anaconda\Lib\site-packages\sklearn\metrics\\_classification.py:1469:  
 UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
[149]: cnn = models.Sequential([
    layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu',
    ↪input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])
```

WARNING:tensorflow:From C:\Users\Pratik  
 Nagare\AppData\Roaming\Python\Python311\site-  
 packages\keras\src\layers\pooling\max\_pooling2d.py:161: The name tf.nn.max\_pool  
 is deprecated. Please use tf.nn.max\_pool2d instead.

```
[150]: cnn.compile(optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy'])
```

```
[151]: cnn.fit(X_train, y_train, epochs=10)
```

Epoch 1/10  
 1563/1563 [=====] - 17s 9ms/step - loss: 2.7309 -  
 accuracy: 0.1037

```
[151]: <keras.src.callbacks.History at 0x1ea900503d0>
```

```
313/313 [=====] - 1s 4ms/step - loss: 2.3026 -
accuracy: 0.1000
```

```
313/313 [=====] - 1s 4ms/step
```

8



```

[0.09971637, 0.10177267, 0.09745668, 0.09848896, 0.10034841,
 0.10028911, 0.10042525, 0.10147531, 0.09965564, 0.10037153],
[0.09971637, 0.10177267, 0.09745668, 0.09848896, 0.10034841,
 0.10028911, 0.10042525, 0.10147531, 0.09965564, 0.10037153]],
dtype=float32)

```

```

[154]: y_classes = [np.argmax(element) for element in y_pred]
       y_classes[:5]

```

```

[154]: [1, 1, 1, 1, 1]

```

```

[156]: y_test[:5]

```

```


[156]: array([3, 8, 8, 0, 6], dtype=uint8)

```

```

[166]: plot_sample(X_test, y_test,3)

```



```

[165]: classes[y_classes[3]]

```

```

[165]: 'automobile'

```

```

[164]: classes[y_classes[3]]

```

```

[164]: 'automobile'

```

```

[ ]:

```

```

[ ]:

```

```

[ ]:

```