

# cnn-cifar10-dataset-1

March 3, 2024

Small Image Classification Using Convolutional Neural Network (CNN)

In this notebook, we will classify small images cifar10 dataset from tensorflow keras datasets. There are total 10 classes as shown below. We will use CNN for classification

```
[1]: import tensorflow as tf
      from tensorflow.keras import datasets, layers, models
      import matplotlib.pyplot as plt
      import numpy as np
```

```
WARNING:tensorflow:From C:\Users\Pratik
Nagare\AppData\Roaming\Python\Python311\site-packages\keras\src\losses.py:2976:
The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use
tf.compat.v1.losses.sparse_softmax_cross_entropy instead.
```

Load the dataset

```
[4]: (X_train, y_train), (X_test, y_test) = datasets.cifar10.load_data()
      X_train.shape
```

```
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 [=====] - 267s 2us/step
```

```
[4]: (50000, 32, 32, 3)
```

```
[5]: X_test.shape
```

```
[5]: (10000, 32, 32, 3)
```

Here we see there are 50000 training images and 10000 test images

```
[6]: y_train.shape
```

```
[6]: (50000, 1)
```

```
[ ]: y_train[:5]
```

y\_train is a 2D array, for our classification having 1D array is good enough. so we will convert this to now 1D array

```
[7]: y_train = y_train.reshape(-1,)
      y_train[:5]
```

```
[7]: array([6, 9, 9, 4, 1], dtype=uint8)
```

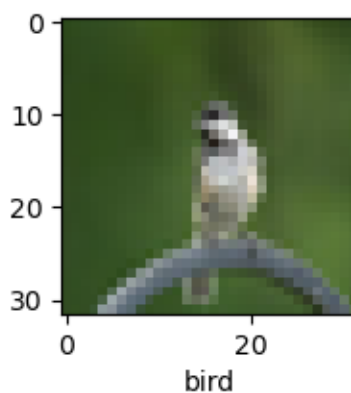
```
[8]: y_test = y_test.reshape(-1,)
```

```
[9]: classes = □
      ↪ ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]
```


Let's plot some images to see what they are

```
[10]: def plot_sample(X, y, index):
        plt.figure(figsize = (15,2))
        plt.imshow(X[index])
        plt.xlabel(classes[y[index]])
```

```
[12]: plot_sample(X_train, y_train, 54)
```



```
[13]: plot_sample(X_train, y_train, 1)
```



Normalize the images to a number from 0 to 1. Image has 3 channels (R,G,B) and each value in the channel can range from 0 to 255. Hence to normalize in 0->1 range, we need to divide it by 255

Normalizing the training data

```
[ ]: X_train = X_train / 255.0
      X_test = X_test / 255.0
```

Build simple artificial neural network for image classification

```
[14]: ann = models.Sequential([
        layers.Flatten(input_shape=(32,32,3)),
        layers.Dense(3000, activation='relu'),
        layers.Dense(1000, activation='relu'),
        layers.Dense(10, activation='sigmoid')
    ])

ann.compile(optimizer='SGD',
            loss='sparse_categorical_crossentropy',
            metrics=['accuracy'])

ann.fit(X_train, y_train, epochs=5)
```

WARNING:tensorflow:From C:\Users\aaive\anaconda3\Lib\site-packages\keras\src\backend.py:873: The name tf.get\_default\_graph is deprecated. Please use tf.compat.v1.get\_default\_graph instead.

WARNING:tensorflow:From C:\Users\aaive\anaconda3\Lib\site-packages\keras\src\optimizers\\_\_init\_\_.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

Epoch 1/5

WARNING:tensorflow:From C:\Users\aaive\anaconda3\Lib\site-packages\keras\src\utils\tf\_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

WARNING:tensorflow:From C:\Users\aaive\anaconda3\Lib\site-packages\keras\src\engine\base\_layer\_utils.py:384: The name tf.executing\_eagerly\_outside\_functions is deprecated. Please use tf.compat.v1.executing\_eagerly\_outside\_functions instead.

986/1563 [=====>...] - ETA: 43s - loss: nan - accuracy: 0.0994

KeyboardInterrupt

Traceback (most recent call last)

Cell In[14], line 12

```
1 ann = models.Sequential([
2     layers.Flatten(input_shape=(32,32,3)),
3     layers.Dense(3000, activation='relu'),
4     layers.Dense(1000, activation='relu'),
5     layers.Dense(10, activation='sigmoid')
6 ])
8 ann.compile(optimizer='SGD',
9             loss='sparse_categorical_crossentropy',
10             metrics=['accuracy'])
---> 12 ann.fit(X_train, y_train, epochs=5)
```

File ~\anaconda3\Lib\site-packages\keras\src\utils\traceback\_utils.py:65, in `filter_traceback.<locals>.error_handler(*args, **kwargs)`

```
63 filtered_tb = None
64 try:
---> 65     return fn(*args, **kwargs)
66 except Exception as e:
67     filtered_tb = _process_traceback_frames(e.__traceback__)
```

File ~\anaconda3\Lib\site-packages\keras\src\engine\training.py:1807, in `Model.fit(self, x, y, batch_size, epochs, verbose, callbacks, validation_split,`

```
validation_data, shuffle, class_weight, sample_weight, initial_epoch,
steps_per_epoch, validation_steps, validation_batch_size, validation_freq,
max_queue_size, workers, use_multiprocessing)
1799 with tf.profiler.experimental.Trace(
1800     "train",
1801     epoch_num=epoch,
1802     ...)
1804     _r=1,
1805 ):
1806     callbacks.on_train_batch_begin(step)
-> 1807     tmp_logs = self.train_function(iterator)
1808     if data_handler.should_sync:
1809         context.async_wait()
```

File ~\anaconda3\Lib\site-packages\tensorflow\python\util\traceback\_utils.py:

```
150, in filter_traceback.<locals>.error_handler(*args, **kwargs)
148 filtered_tb = None
149 try:
--> 150     return fn(*args, **kwargs)
151 except Exception as e:
152     filtered_tb = _process_traceback_frames(e.__traceback__)
```

File

```
~\anaconda3\Lib\site-packages\tensorflow\python\eager\polymorphic_function\polymorphic_function.py:832, in Function.__call__(self, *args, **kws)
829 compiler = "xla" if self._jit_compile else "nonXla"
```

```

    831 with OptionalXlaContext(self._jit_compile):
--> 832     result = self._call(*args, **kwds)
    834 new_tracing_count = self.experimental_get_tracing_count()
    835 without_tracing = (tracing_count == new_tracing_count)

File
↳ ~\anaconda3\Lib\site-packages\tensorflow\python\eager\polymorphic_function\polymorphic_function.py:868, in Function._call(self, *args, **kwds)
    865     self._lock.release()
    866     # In this case we have created variables on the first call, so we run
↳ the
    867     # defunned version which is guaranteed to never create variables.
--> 868     return tracing_compilation.call_function(
    869         args, kwds, self._no_variable_creation_config
    870     )
    871 elif self._variable_creation_config is not None:
    872     # Release the lock early so that multiple threads can perform the call
    873     # in parallel.
    874     self._lock.release()

File
↳ ~\anaconda3\Lib\site-packages\tensorflow\python\eager\polymorphic_function\tracing_compilation.py:139, in call_function(args, kwargs, tracing_options)
    137 bound_args = function.function_type.bind(*args, **kwargs)
    138 flat_inputs = function.function_type.unpack_inputs(bound_args)
--> 139 return function._call_flat( # pylint: disable=protected-access
    140     flat_inputs, captured_inputs=function.captured_inputs
    141 )

File
↳ ~\anaconda3\Lib\site-packages\tensorflow\python\eager\polymorphic_function\concrete_function.py:1323, in ConcreteFunction._call_flat(self, tensor_inputs, captured_inputs)
    1319 possible_gradient_type = gradients_util.PossibleTapeGradientTypes(args)
    1320 if (possible_gradient_type == gradients_util.POSSIBLE_GRADIENT_TYPES_NO_TAPES
    1321     and executing_eagerly):
    1322     # No tape is watching; skip to running the function.
-> 1323     return self._inference_function.call_preflattened(args)
    1324 forward_backward = self._select_forward_and_backward_functions(
    1325     args,
    1326     possible_gradient_type,
    1327     executing_eagerly)
    1328 forward_function, args_with_tangents = forward_backward.forward()

File
↳ ~\anaconda3\Lib\site-packages\tensorflow\python\eager\polymorphic_function\atomic_function.py:216, in AtomicFunction.call_preflattened(self, args)
    214 def call_preflattened(self, args: Sequence[core.Tensor]) -> Any:
    215     """Calls with flattened tensor inputs and returns the structured
↳ output."""

```

```

--> 216 flat_outputs = self.call_flat(*args)
      217 return self.function_type.pack_output(flat_outputs)

```

File

```

-> ~\anaconda3\Lib\site-packages\tensorflow\python\eager\polymorphic_function\atomic_function
py:251, in AtomicFunction.call_flat(self, *args)
      249 with record.stop_recording():
      250     if self._bound_context.executing_eagerly():
--> 251         outputs = self._bound_context.call_function(
      252             self.name,
      253             list(args),
      254             len(self.function_type.flat_outputs),
      255         )
      256     else:
      257         outputs = make_call_op_in_graph(
      258             self,
      259             list(args),
      260             self._bound_context.function_call_options.as_attrs(),
      261         )

```

File ~\anaconda3\Lib\site-packages\tensorflow\python\eager\context.py:1486, in

```

-> Context.call_function(self, name, tensor_inputs, num_outputs)
      1484 cancellation_context = cancellation.context()
      1485 if cancellation_context is None:
-> 1486     outputs = execute.execute(
      1487         name.decode("utf-8"),
      1488         num_outputs=num_outputs,
      1489         inputs=tensor_inputs,
      1490         attrs=attrs,
      1491         ctx=self,
      1492     )
      1493 else:
      1494     outputs = execute.execute_with_cancellation(
      1495         name.decode("utf-8"),
      1496         num_outputs=num_outputs,
      (...)
      1500         cancellation_manager=cancellation_context,
      1501     )

```

File ~\anaconda3\Lib\site-packages\tensorflow\python\eager\execute.py:53, in

```

-> quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
      51 try:
      52     ctx.ensure_initialized()
----> 53     tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_name,
      54                                         inputs, attrs, num_outputs)
      55 except core._NotOkStatusException as e:
      56     if name is not None:

```

KeyboardInterrupt:

You can see that at the end of 5 epochs, accuracy is at around 48.48%

```
[ ]: from sklearn.metrics import confusion_matrix , classification_report
import numpy as np
y_pred = ann.predict(X_test)
y_pred_classes = [np.argmax(element) for element in y_pred]

print("Classification Report: \n", classification_report(y_test,
↪y_pred_classes))
```

Now let us build a convolutional neural network to train our images

```
[ ]: cnn = models.Sequential([
    layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu',
↪input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])
```

```
[ ]: cnn.compile(optimizer='adam',
                loss='sparse_categorical_crossentropy',
                metrics=['accuracy'])
```

```
[ ]: cnn.fit(X_train, y_train, epochs=10)
```

With CNN, at the end 5 epochs, accuracy was at around 70.28% which is a significant improvement over ANN. CNN's are best for image classification and gives superb accuracy. Also computation is much less compared to simple ANN as maxpooling reduces the image dimensions while still preserving the features

```
[ ]: cnn.evaluate(X_test,y_test)
```

```
[ ]: y_pred = cnn.predict(X_test)
y_pred[:5]
```

```
[ ]: y_classes = [np.argmax(element) for element in y_pred]
y_classes[:5]
```

```
[ ]: y_test[:5]
```

```
[ ]: plot_sample(X_test, y_test,3)
```

```
[ ]: classes[y_classes[3]]
```

```
[ ]: classes[y_classes[3]]
```

### Exercise

Use CNN to do handwritten digits classification using MNIST dataset. You can use this notebook as a reference: [https://github.com/codebasics/py/blob/master/DeepLearningML/1\\_digits\\_recognition/digits\\_recog](https://github.com/codebasics/py/blob/master/DeepLearningML/1_digits_recognition/digits_recog)

Above we used ANN for digits classification. You need to modify this code to use CNN instead. Check how accuracy improves fast with CNN and figure out how CNN can be a better choice for doing image classification compared to ANN. Once you have worked on this problem on your own, you can check my solution by clicking on this link: [Solution](#)